

# HomeWork

December 9, 2018

```
In [1]: from sqlalchemy import create_engine
        from sqlalchemy.exc import ResourceClosedError
        from sqlalchemy.types import VARCHAR
        from functools import partial
        import pandas as pd
        from scipy.spatial import distance
        from sklearn.neighbors import NearestNeighbors
        from IPython.display import display, HTML, Markdown, Latex
        from sklearn.utils.extmath import randomized_svd
        import numpy as np
```

## 0.1 Reading the data from database into a pandas dataframe

```
In [2]: #connection
        def DatabaseConnect(username, password, schema):
            conn_str = "mysql+pymysql://{username}:{password}@localhost/{schema}?charset=utf8&"
                        .format(username=username, password=password, schema=schema)
            engine = create_engine(conn_str, pool_recycle=1800)
            return engine
        RecSysConnect = partial(DatabaseConnect, 'root', 'mysql-password', 'recsys')
        e = RecSysConnect()
```

```
In [3]: #read ratings
        sql_cmt = "select userId, movieId, rating from ml100k_ratings;"
        raw_rating = pd.read_sql(sql_cmt, con=e)
        rating = raw_rating.pivot(index="userId", columns="movieId", values="rating")
        rating.head(10)
```

```
Out[3]: movieId  1      2      3      4      5      6      7      8      9     10     ...  \
userId
1           5.0    3.0    4.0    3.0    3.0    5.0    4.0    1.0    5.0    3.0    ...
2           4.0    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    2.0    ...
3           NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    ...
4           NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    ...
5           4.0    3.0    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    ...
6           4.0    NaN    NaN    NaN    NaN    NaN    NaN    2.0    4.0    4.0    NaN    ...
7           NaN    NaN    NaN    5.0    NaN    NaN    NaN    5.0    5.0    5.0    4.0    ...
8           NaN    NaN    NaN    NaN    NaN    NaN    NaN    3.0    NaN    NaN    NaN    ...
```

9	NaN	NaN	NaN	NaN	NaN	5.0	4.0	NaN	NaN	NaN	...
10	4.0	NaN	NaN	4.0	NaN	NaN	4.0	NaN	4.0	NaN	...

movieId	1673	1674	1675	1676	1677	1678	1679	1680	1681	1682
userId										
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
5	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
6	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
7	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
8	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
10	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

[10 rows x 1682 columns]

In [4]: *#read holdout data*

```
sql_cmt = "select userId from ml100k_ratings group by userId order by count(*) DESC limit 100"
user_100 = pd.read_sql(sql_cmt, con=e)['userId']
sql_cmt = "select movieId from ml100k_ratings group by movieId order by count(*) DESC limit 100"
movie_10 = pd.read_sql(sql_cmt, con=e)['movieId']
holdout = rating.loc[user_100,movie_10]
```

In [5]: rating.loc[user\_100,movie\_10]=np.nan

## 0.2 Calculation of similarities

In [6]: `import sklearn.metrics.pairwise`

```
def center(df):
    return df.sub( df.mean(axis=1), axis=0 )

def cosine(df, axis=0):
    dff = df.fillna(0)
    if axis == 0: # Columns
        return pd.DataFrame(sklearn.metrics.pairwise.cosine_similarity(dff.T), index=dff.index, columns=dff.columns)
    else:
        return pd.DataFrame(sklearn.metrics.pairwise.cosine_similarity(dff), index=dff.index, columns=dff.columns)
```

```
In [7]: r_cent = center(rating)
userSim = rating.T.corr(method='pearson')
#userSim = cosine(r_cent, axis=1)
itemSim = cosine(r_cent)
display(userSim.head(),itemSim.head())
```

userId	1	2	3	4	5	6	7	\
userId								

1	1.000000	0.156868	-0.069171	-0.666667	0.386678	0.298236	0.260460
2	0.156868	1.000000	0.067420	0.148522	0.327327	0.446269	0.381875
3	-0.069171	0.067420	1.000000	-0.262600	NaN	-0.109109	0.104828
4	-0.666667	0.148522	-0.262600	1.000000	1.000000	-0.581318	-0.660529
5	0.386678	0.327327	NaN	1.000000	1.000000	0.241817	0.170291

userId	8	9	10	...	934	935	936 \
userId				...			
1	0.694108	-0.301511	-0.124713	...	0.066802	-0.252820	0.435300
2	0.585491	0.242536	0.668145	...	0.021007	-0.271163	0.214017
3	0.291937	NaN	0.311086	...	NaN	NaN	-0.045162
4	0.642938	NaN	-0.301511	...	0.500000	NaN	-0.203653
5	0.537400	0.577350	0.087343	...	0.229532	-0.500000	0.439286

userId	937	938	939	940	941	942	943
userId							
1	-0.044483	0.157353	0.584437	0.257908	0.000000	-0.202777	0.048947
2	0.561645	0.331587	0.000000	-0.011682	-0.062017	0.085960	0.479702
3	0.000000	-0.137523	NaN	-0.104678	1.000000	-0.011792	NaN
4	NaN	0.375000	NaN	0.850992	1.000000	0.412568	NaN
5	0.608581	0.484211	0.880705	0.027038	0.468521	0.318163	0.346234

[5 rows x 943 columns]

movieId	1	2	3	4	5	6	7 \
movieId							
1	1.000000	-0.023667	-0.042012	0.012120	0.007037	-0.002027	0.020884
2	-0.023667	1.000000	0.031568	0.030879	0.013452	-0.010600	-0.053574
3	-0.042012	0.031568	1.000000	-0.118348	0.012564	0.054819	-0.087453
4	0.012120	0.030879	-0.118348	1.000000	-0.149492	-0.022921	0.005652
5	0.007037	0.013452	0.012564	-0.149492	1.000000	-0.041295	-0.043651

movieId	8	9	10	...	1673	1674	1675 \
movieId				...			
1	0.100244	-0.046859	0.002856	...	0.066344	0.0	0.000000
2	-0.007358	-0.103660	-0.020773	...	0.000000	0.0	0.000000
3	-0.151500	-0.066376	-0.049901	...	0.000000	0.0	0.000000
4	0.111663	0.029485	0.009932	...	0.000000	0.0	-0.114413
5	0.012695	-0.061202	-0.033263	...	0.000000	0.0	0.000000

movieId	1676	1677	1678	1679	1680	1681	1682
movieId							
1	0.000000	0.012261	0.0	0.0	0.0	0.000000	0.000000
2	0.000000	0.000000	0.0	0.0	0.0	0.003661	0.034941
3	0.000000	0.201111	0.0	0.0	0.0	0.000000	0.031866
4	-0.114413	0.090004	0.0	0.0	0.0	0.002727	-0.048234
5	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.043673

[5 rows x 1682 columns]

## 0.3 Recommenders

### 0.3.1 - randomized\_svd

```
In [8]: def svd_pred(rating,holdout):
        rating.columns = rating.columns.map(str)
        means = rating.mean(axis=1)
        cent = rating.sub( means, axis=0 )
        U_, Sigma, VT_ = randomized_svd(cent.to_sparse().to_coo().tocsc(), n_components=2,
        U = pd.DataFrame(U_, index=rating.index)
        VT = pd.DataFrame(VT_, columns=rating.columns)
        full = U.mul(Sigma).dot(VT)
        full.columns=full.columns.map(int)
        means = rating.mean(axis=1).loc[user_100]
        prediction = full.loc[user_100,movie_10].add(means,axis=0)
        y_true=[]
        y_predicted=[]
        for u in holdout.index:
            for m in holdout.columns:
                if (holdout.isna().loc[u,m]==False):
                    y_true=y_true+[holdout.loc[u,m]]
                    y_predicted=y_predicted+[prediction.loc[u,m]]
        return y_true, y_predicted
```

```
In [9]: y_true_svd, y_pred_svd = svd_pred(rating,holdout)
```

```
In [10]: pd.DataFrame([y_true_svd,y_pred_svd])
```

```
Out[10]:
```

	0	1	2	3	4	5	6	\
0	5.000000	5.000000	5.000000	4.000000	2.00000	3.000000	3.000000	
1	3.969673	3.624942	2.092731	3.329952	2.95191	3.541682	2.897582	
	7	8	9	...	745	746	747	\
0	3.000000	3.000000	3.000000	...	5.000000	4.000000	5.00000	
1	2.428885	2.977192	2.705344	...	3.488522	3.394138	3.49022	
	748	749	750	751	752	753	754	
0	5.000000	5.00000	4.000000	4.000000	4.00000	4.000000	1.000000	
1	3.973007	3.68861	4.086564	3.667382	3.31407	3.672904	3.285864	

[2 rows x 755 columns]

### 0.3.2 - user-based

```
In [11]: def recommendUser(df, user, item, k=3, weighted=True):
        knn = userSim.loc[user].drop(user).nlargest(k)
```

```

norm = knn.sum()
rating = 0
for u in knn.index:
    rating += df.loc[u, item] * knn[u]
rating = np.clip( rating / norm, 1, 5 )
return (rating)

```

```

In [12]: def user_based(rating,holdout):
rating.columns = rating.columns.map(int)
y_true=[]
y_predicted=[]
for u in holdout.index:
    for i in holdout.columns:
        if (holdout.isna().loc[u,i]==False):
            y_true=y_true+[holdout.loc[u,i]]
            y_predicted=y_predicted+[recommendUser(rating, u, i, 2)]
    return y_true,y_predicted
y_true_user, y_pred_user = user_based(rating,holdout)
df=pd.DataFrame([y_true_user,y_pred_user]).dropna(axis=1)
y_true_user = df.iloc[0]
y_pred_user = df.iloc[1]
display(df)

```

	2	7	9	11	17	18	21	27	28	46	\
0	5.0	3.000000	3.0	3.0	2.000000	3.000000	1.000000	4.0	4.0	4.000000	
1	3.0	3.989795	4.0	4.0	2.946724	3.579914	3.473362	3.0	4.0	4.463657	
...	683	688	699	707	717	718	720	722	734	752	
0	...	5.0	3.0	5.0	5.0	3.0	5.0	3.0	4.0	3.0	4.0
1	...	3.5	2.0	3.0	3.5	3.0	3.0	2.5	4.5	3.5	2.5

[2 rows x 104 columns]

### 0.3.3 - item-based

```

In [13]: def recommendItem(df,user, item, k=2, weighted=True):
knn = itemSim.loc[item][ df.loc[user].notnull() ]
knn = knn[ knn> 0 ].drop(item, errors='ignore').nlargest(k)
if weighted:
    norm = knn.sum()
else:
    knn.values.fill(1/len(knn))
    norm = 1
ratings = 0
for i in knn.index:
    ratings += df.loc[user,i] * knn[i]
ratings = np.clip( ratings / norm, 1, 5 )

```

```

    return (ratings)
def item_based(rating,holdout):
    rating.columns = rating.columns.map(int)
    y_true=[]
    y_predicted=[]
    for u in holdout.index:
        for i in holdout.columns:
            if (holdout.isna().loc[u,i]==False):
                y_true=y_true+[holdout.loc[u,i]]
                y_predicted=y_predicted+[recommendItem(rating, u, i)]
    return y_true,y_predicted
y_true_item, y_pred_item = item_based(rating,holdout)
pd.DataFrame([y_true_item,y_pred_item])

```

```

Out[13]:
   0    1    2    3    4    5    6    7    8  \
0  5.0  5.0  5.000000  4.000000  2.0  3.000000  3.000000  3.000000  3.000000
1  5.0  5.0  2.857781  3.531249  4.0  4.048189  3.601238  2.499696  2.429392

      9    ...  745  746    747    748    749    750  \
0  3.000000  ...  5.0  4.0  5.000000  5.000000  5.000000  4.000000
1  2.587803  ...  4.0  4.0  4.417514  3.531249  3.490975  4.551081

      751    752    753    754
0  4.000000  4.000000  4.000000  1.000000
1  3.601238  3.473194  3.468862  2.835027

[2 rows x 755 columns]

```

## 0.4 Evaluation: MSE MAE

```

In [14]: from sklearn.metrics import mean_squared_error,mean_absolute_error
def evaluation(y_true, y_predicted):
    result=pd.DataFrame([y_true,y_predicted])
    print(result)
    #MAE=np.mean(np.abs(result[0] - result[1]))/755
    MAE = mean_absolute_error(y_true,y_predicted)
    RMSE = np.sqrt(mean_squared_error(y_true,y_predicted))
    #RMSE=np.mean((result[0] - result[1])**2)/755
    display("MAE",MAE,"RMSE",RMSE)

```

```

In [15]: evaluation(y_true_svd, y_pred_svd)

```

```

      0    1    2    3    4    5    6  \
0  5.000000  5.000000  5.000000  4.000000  2.000000  3.000000  3.000000
1  3.969673  3.624942  2.092731  3.329952  2.95191  3.541682  2.897582

      7    8    9    ...  745    746    747  \
0  3.000000  3.000000  3.000000  ...  5.000000  4.000000  5.000000
1  2.428885  2.977192  2.705344  ...  3.488522  3.394138  3.49022

```

	748	749	750	751	752	753	754
0	5.000000	5.000000	4.000000	4.000000	4.000000	4.000000	1.000000
1	3.973007	3.68861	4.086564	3.667382	3.31407	3.672904	3.285864

[2 rows x 755 columns]

'MAE'

0.7386856980861768

'RMSE'

0.9117497430960818

In [16]: evaluation(y\_true\_item, y\_pred\_item)

	0	1	2	3	4	5	6	7	8	\
0	5.0	5.0	5.000000	4.000000	2.0	3.000000	3.000000	3.000000	3.000000	
1	5.0	5.0	2.857781	3.531249	4.0	4.048189	3.601238	2.499696	2.429392	

  

	9	...	745	746	747	748	749	750	\
0	3.000000	...	5.0	4.0	5.000000	5.000000	5.000000	4.000000	
1	2.587803	...	4.0	4.0	4.417514	3.531249	3.490975	4.551081	

  

	751	752	753	754
0	4.000000	4.000000	4.000000	1.000000
1	3.601238	3.473194	3.468862	2.835027

[2 rows x 755 columns]

'MAE'

0.6783970962926593

'RMSE'

0.9412219376564158

In [17]: evaluation(y\_true\_user, y\_pred\_user)

	2	7	9	11	17	18	21	27	28	46	\
0	5.0	3.000000	3.0	3.0	2.000000	3.000000	1.000000	4.0	4.0	4.000000	
1	3.0	3.989795	4.0	4.0	2.946724	3.579914	3.473362	3.0	4.0	4.463657	
...	683	688	699	707	717	718	720	722	734	752	
0	...	5.0	3.0	5.0	5.0	3.0	5.0	3.0	4.0	3.0	4.0
1	...	3.5	2.0	3.0	3.5	3.0	3.0	2.5	4.5	3.5	2.5

[2 rows x 104 columns]

'MAE'

1.0512574811196467

'RMSE'

1.2742844501470791

---

```
In [14]: #Use surprise package
from surprise import SVD
from surprise.model_selection import cross_validate
from surprise import Reader, Dataset
reader = Reader(rating_scale=(1, 5))
# The columns must correspond to user id, item id and ratings (in that order).
data = Dataset.load_from_df(raw_rating, reader)
cross_validate(SVD(), data, measures=['RMSE', 'MAE'], cv=2, verbose=True)
```

Evaluating RMSE, MAE of algorithm SVD on 2 split(s).

	Fold 1	Fold 2	Mean	Std
RMSE (testset)	0.9598	0.9560	0.9579	0.0019
MAE (testset)	0.7583	0.7563	0.7573	0.0010
Fit time	3.38	4.16	3.77	0.39
Test time	0.66	0.59	0.63	0.04

```
Out[14]: {'fit_time': (3.375916004180908, 4.163803815841675),
          'test_mae': array([0.75829125, 0.7563168 ]),
          'test_rmse': array([0.95982148, 0.95600433]),
          'test_time': (0.6646459102630615, 0.5870378017425537)}
```